

2019 Git Contributor Summit – Topics

Elijah Newren

Palantir Technologies

2019-02-05

Setup

- ▶ 30-40 people
 - ▶ I got to put lots of faces to names
 - ▶ Had many top contributors, but definitely not all of them
- ▶ Big projects & companies represented:
 - ▶ GitHub, GitLab, Atlassian, Gerrit, Microsoft, Google
- ▶ Peff's comment on merge work

Overview

- ▶ Fetch response & CDN offloading
- ▶ A review of "Big"
- ▶ NewHash
- ▶ Poll: prove vs. jumbled output
- ▶ Poll: where should Git Merge be next year?
- ▶ Using mailmap by default in git log?
- ▶ Fetch response sideband-all
- ▶ Protocol v2 for push
- ▶ Partial Clones
- ▶ Informal Break discussions
- ▶ Breakout Sessions
- ▶ CI

Fetch response & CDN offloading

- ▶ allows resumable cloning
- ▶ does load balancing
- ▶ gets the static part of history (e.g. until a week ago) from CDN, and last bits from "main" server
- ▶ questions about whether to do multiple bits offloaded (e.g. almost full clone, only stuff from last month, etc.); can server keep track of manifest and direct client to necessary subset of pack on a CDN?

A review of "Big"

- ▶ references, history, wide-checkout, deep-checkout, lots to gc, etc.
- ▶ newer stuff: partial clones, worktrees, commit-graph
- ▶ originally planned to do a breakout session later, but...
- ▶ git-sizer (usage by GitLab vs. GitHub)
- ▶ large blobs, partial clones
- ▶ partial or hierarchical indexes

NewHash

- ▶ sha1 -> sha256
- ▶ have sha256 repo locally talking to a server using sha1?
- ▶ as of yesterday, binary that can create either sha1 or sha256 repos
- ▶ will be using fixed length listing of shas in packfile; if given sha1 is fourth in list, then the corresponding sha256 will be fourth
- ▶ next: interoperation; fetch & push coming up next
- ▶ done a fair amount of work so moving to a new hash in the future with a different length should be much less work
- ▶ no automatic translation of commit messages, but maintenance of dual-mapping of hashes
- ▶ (Comments on sha1dc & its performance)
- ▶ Submodules is the biggest issue right now

Poll: prove vs. jumbled output

Poll: where should Git Merge be next year?

Using mailmap by default in git log?

- ▶ People change names for lots of reasons (including transliteration differences)
- ▶ Keep an option to not use mailmap
- ▶ People generally positive on the idea

Fetch response sideband-all

- ▶ sidebands for progress messages and errors
- ▶ sideband currently limited to when sending packfile
- ▶ proposal: expand sideband for whole response, not just packfile.
- ▶ particularly useful given ideas to do CDN
- ▶ also needed for keep-alive messages
- ▶ this will be a negotiated new capability (can't do it backward-compatibly)

Protocol v2 for push

- ▶ ref advertisement the main issue
- ▶ Some want to modify the commit server-side (?!?!?)
 - ▶ tweak the commit message somehow
 - ▶ rebase-on-push
 - ▶ reformat-on-push
- ▶ discussion of how to split messages up into sub-commands
- ▶ a way to retry pushes without re-pushing everything (e.g. someone else updated the branch, you then re-merged or rebased locally and want to push again, meaning the server already has *most* the objects but just needs a few new ones)

Partial Clones

- ▶ doing work to have multiple remotes (related: CDN usage)
- ▶ still very tied to having a server around to request additional objects
- ▶ we need to have a way to keep upload-pack open and do multiple requests
- ▶ has rudimentary ability to filter trees
 - ▶ Matthew Devore doing some work in this area right now, but it appears to be based on depth rather than width?
- ▶ connection with sparse checkout is kind of hacky right now
- ▶ there are reachability enforcement issues in V2, which becomes even more of an issue with partial clones (now need to worry about blobs not just commits)
- ▶ in a partial clone world, server can't gc
- ▶ sidenote: dumb http support
 - ▶ no major hosting provider supports it
 - ▶ some people like it due to resumability (e.g. Joey Hess & git annex)
 - ▶ cgit (repo browser) provides dumb http support natively

CI

- ▶ Folks aren't all that happy with the Travis setup
- ▶ Dscho has a lot of machinery built up around Azure Pipelines
- ▶ PRs to `github.com:gitgitgadget/git` will automatically be built on Windows, MacOS X, and linux
- ▶ Interest in getting emails for failures that their topic branch caused (note: get topic author from tip commit author if not Junio)
- ▶ This may be able to move to `github.com:git/git` after Dscho's patches merge down

Breakout Sessions

Four simultaneous breakout sessions were held:

- ▶ merging
- ▶ GSoC
- ▶ structured logging
- ▶ windows big files

Breakout Session: merging

I led the breakout session on merging

- ▶ merge-recursive rewrite
 - ▶ questions and basic explanation of how the algorithm works
 - ▶ want incremental updates on merge-recursive rewrite
 - ▶ make merge-recursive code part of libgit.a ?
 - ▶ people are very happy about idea to not touch the working tree
- ▶ make rebase --merge the default
 - ▶ use performance tests to see how well it compares (p3400-rebase.sh)
 - ▶ may later also reimplement the am-specific flags on top of interactive
- ▶ make use of best merge bases in more places (e.g. git diff A...B uses a suboptimal one)
- ▶ rebase --rebase-merges; more thorough reimplementation ideas

Informal Break discussions

- ▶ saving incomplete merge state
- ▶ checkout overlay
- ▶ filter-repo