

Scaling the Merge Machinery

Elijah Newren

Palantir Technologies

Git Merge

Elijah Newren

Palantir Technologies

(Git Merge)²

Elijah Newren

Palantir Technologies

Outline

1 “Merge machinery”

- Affected commands
- The journey
- Overhaul
- Types of performance strategies

2 Merging and renames background

3 Strategies to improve rename performance

4 Results

Affected Commands

The merge machinery (merge-recursive) powers several aspects of git:

- merge
- cherry-pick
- revert
- rebase
- am -3
- stash
- checkout -m

The journey

A few years ago...

Issues starting the journey

- cherry-pick would fail to detect renames and fail to notify about needed `merge.renameLimit`
- cherry-pick would ignore `merge.renameLimit > 32767`
- if directory renames involved, files would be left in wrong directory
- people wrote custom purpose scripts to cherry-pick things
- after fixing `merge.renameLimit`, cherry-picking small patches would take more than 9 minutes.

Quotes

The two most prolific authors of git opining on merge-recursive:

Quotes

The two most prolific authors of git opining on merge-recursive:

- “[It is] some pretty hairy code. Every time I start to look at it I get confused and can’t remember what breakthrough I thought I was close to making before.” (Jeff King)

Quotes

The two most prolific authors of git opining on merge-recursive:

- “[It is] some pretty hairy code. Every time I start to look at it I get confused and can’t remember what breakthrough I thought I was close to making before.” (Jeff King)
- “I’ve written off that code as mostly unsalvageable long time ago.” (Junio Hamano)

Goals

Goals for my rewrite of the machinery are to improve each of:

- Maintainability & understandability
- API Quality (enable new features?)
- Correctness
- Performance

Goals

Goals for my rewrite of the machinery are to improve each of:

- Maintainability & understandability
- API Quality (enable new features?)
- Correctness
- **Performance**

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

- Squeezing performance out of the hardware

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

- Squeezing performance out of the hardware
- Applying ideas from other problem domains to new areas

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

- Squeezing performance out of the hardware
- Applying ideas from other problem domains to new areas
- Using clever approximation algorithms to get near solutions

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

- Squeezing performance out of the hardware
- Applying ideas from other problem domains to new areas
- Using clever approximation algorithms to get near solutions
- Inventing new algorithms

Types of performance strategies

I have always enjoyed performance talks; they make me feel smarter:

- Squeezing performance out of the hardware
- Applying ideas from other problem domains to new areas
- Using clever approximation algorithms to get near solutions
- Inventing new algorithms

Types of performance strategies

Actual performance strategies used:

Types of performance strategies

Actual performance strategies used:

- Don't do unnecessary work

Types of performance strategies

Actual performance strategies used:

- Don't do unnecessary work
- Don't redo work

Types of performance strategies

Actual performance strategies used:

- Don't do unnecessary work
- Don't redo work
- Don't redo unnecessary work

Types of performance strategies

Actual performance strategies used:

- Don't do unnecessary work
- Don't redo work
- Don't redo unnecessary work
- Fudge "unnecessary"

Warning

- Glossing over lots of details
- Simplifications not fully accurate

Outline

- 1 “Merge machinery”
- 2 Merging and renames background
 - Content merge
 - Combining content merges
 - Why renames are important
 - How rename detection works
- 3 Strategies to improve rename performance
- 4 Results

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
?????  
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
?????  
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
?????  
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
?????  
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
?????  
shiver(me.timbers);
```


Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
explore_sea(plus, plus);  
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...
speak_like_a_pirate(arrrgs);
explore_sea(aye, matey);
shiver(me.timbers);
...
```

Same file from branch Side2:

```
...
speak_like_a_pirate(arrrgs);
explore_sea(me.love[0]);
shiver(me.timbers);
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);
explore_sea(plus, plus);
shiver(me.timbers);
```

Which results in the following merge:

```
speak_like_a_pirate(arrrgs);
<<<<<< HEAD
explore_sea(aye, matey);
=====
explore_sea(me.love[0]);
>>>>>> branchB
shiver(me.timbers);
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
explore_sea(plus, plus);  
shiver(me.timbers);
```

Shorthand:

```
path  
Base : hash_orig  
Side1: hash_A  
Side2: hash_B
```

Example:

```
buccaneer.c  
Base : ba771ed  
Side1: 57abbed  
Side2: b1a57ed
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
explore_sea(plus, plus);  
shiver(me.timbers);
```

Shorthand:

```
path  
Base : hash_orig  
Side1: hash_A  
Side2: hash_B
```

Example:

```
buccaneer.c  
Base : ba771ed  
Side1: 57abbed  
Side2: b1a57ed
```

Three-way content merge

File from branch Side1:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(aye, matey);  
shiver(me.timbers);  
...
```

Same file from branch Side2:

```
...  
speak_like_a_pirate(arrrgs);  
explore_sea(me.love[0]);  
shiver(me.timbers);  
...
```

Correct merge depends on the version in the merge base:

```
speak_like_a_pirate(arrrgs);  
explore_sea(plus, plus);  
shiver(me.timbers);
```

Shorthand:

```
path  
Base : hash_orig  
Side1: hash_A  
Side2: hash_B
```

Example:

```
buccaneer.c  
Base : ba771ed  
Side1: 57abbed  
Side2: b1a57ed
```

Note: If any two of the hashes match, we can resolve without looking at the contents of the file.

Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge

Merge

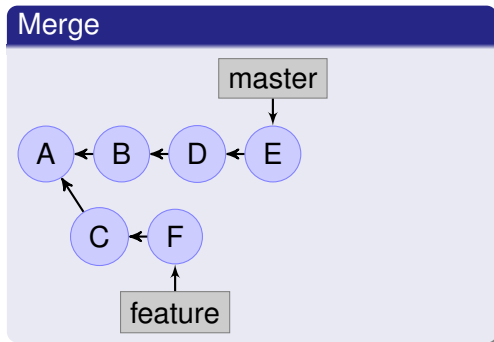
Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge



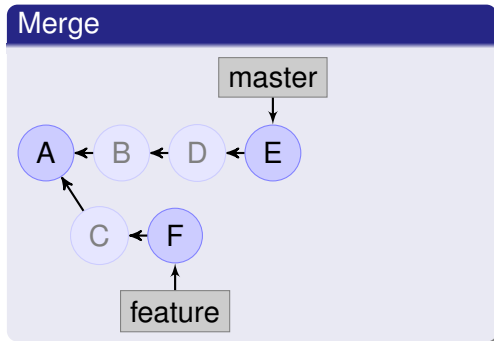
Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge




Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge



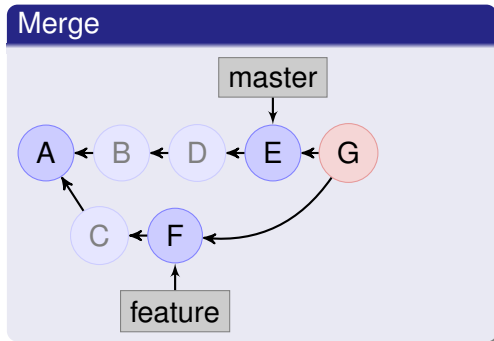
Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge



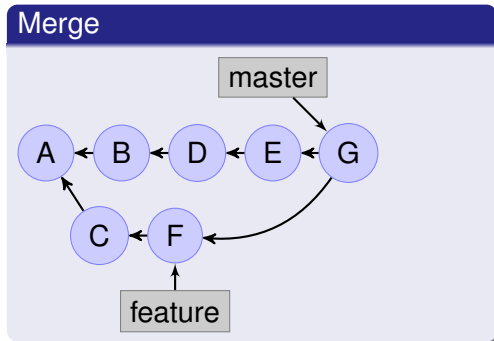
Three-way Merging

```
$ git checkout master
```

```
$ git merge feature
```

Get three relevant trees, then for each path:

- Get version of path in each tree
- Do three-way content merge




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```


Starting state



Three-way Merging

```
$ git checkout master  
$ git cherry-pick C..feature
```


Picking F



Three-way Merging

```
$ git checkout master  
$ git cherry-pick C..feature
```

Picking F




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picking F




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picked F




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picking G




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picking G




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picking G




Three-way Merging

```
$ git checkout master
```

```
$ git cherry-pick C..feature
```

Picked G




Three-way Merging

```
$ git checkout master
```

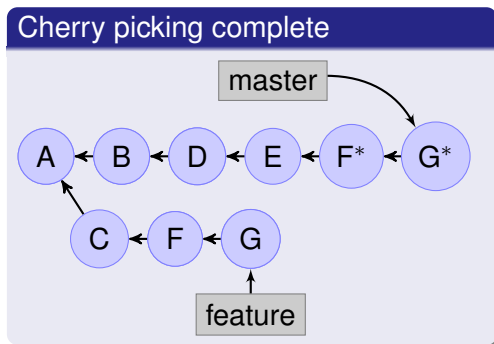
```
$ git cherry-pick C..feature
```

Cherry picking complete



Three-way Merging

```
$ git checkout master  
$ git cherry-pick C..feature
```



Rebasing and reverting are handled similarly to cherry-picking.

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

Changes to be committed:

new file: viking.c

Unmerged paths:

deleted by them: buccaneer.c

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- **buccaneer.c: modify/delete conflict**
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c

Unmerged paths:
  deleted by them:   buccaneer.c
```


Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- **buccaneer.c: modify/delete conflict**
- **viking.c: totally new file**
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c

Unmerged paths:
  deleted by them:   buccaneer.c
```

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

Changes to be committed:

new file: viking.c

Unmerged paths:

deleted by them: buccaneer.c

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

Changes to be committed:

new file: viking.c

Unmerged paths:

deleted by them: buccaneer.c

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	ba771e5
Side1:	e5ca185
Side2:	defea75
Merged:	acc0575

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them: buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	ba771e5
Side1:	e5ca185
Side2:	defea75
Merged:	acc0575

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified: viking.c
```

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	ba771e5
Side1:	e5ca185
Side2:	defea75
Merged:	acc0575

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them: buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	ba771e5
Side1:	e5ca185
Side2:	defea75
Merged:	acc0575

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified: viking.c
```

Why renames are important

If a rename is not detected:

	buccaneer.c	viking.c
Base:	ba771e5	0000000
Side1:	e5ca185	0000000
Side2:	0000000	defea75

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	ba771e5
Side1:	e5ca185
Side2:	defea75
Merged:	acc0575

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```


How rename detection works

How does git detect renames? For each side...

Files in Base	Files in given side
README.md	README.md
archery.js	corrupt.js
baseball.js	divine.js
build.log	dull.js
football.js	grand.js
golf.js	lame.js
running.js	

For each pair of files, what percentage of lines are found in both?

	corrupt.js	divine.js	dull.js	grand.js	lame.js
archery.js					
baseball.js					
build.log					
football.js					
golf.js					
running.js					

Matrix of similarity percentages

How rename detection works

How does git detect renames? For each side...

Files in Base	Files in given side
README.md	README.md
archery.js	corrupt.js
baseball.js	divine.js
build.log	dull.js
football.js	grand.js
golf.js	lame.js
running.js	

For each pair of files, what percentage of lines are found in both?

	corrupt.js	divine.js	dull.js	grand.js	lame.js
archery.js					
baseball.js					
build.log					
football.js					
golf.js					
running.js					

Matrix of similarity percentages

How rename detection works

How does git detect renames? For each side...

Files in Base	Files in given side
README.md	README.md
archery.js	corrupt.js
baseball.js	divine.js
build.log	dull.js
football.js	grand.js
golf.js	lame.js
running.js	

For each pair of files, what percentage of lines are found in both?

	corrupt.js	divine.js	dull.js	grand.js	lame.js
archery.js					
baseball.js					
build.log					
football.js					
golf.js					
running.js					

Matrix of similarity percentages

How rename detection works

How does git detect renames? For each side...

Files in Base	Files in given side
README.md	README.md
archery.js	corrupt.js
baseball.js	divine.js
build.log	dull.js
football.js	grand.js
golf.js	lame.js
running.js	

For each pair of files, what percentage of lines are found in both?

	corrupt.js	divine.js	dull.js	grand.js	lame.js
archery.js					
baseball.js					
build.log					
football.js					
golf.js					
running.js					

Matrix of similarity percentages

How rename detection works

How does git detect renames? For each side...

Files in Base	Files in given side
README.md	README.md
archery.js	corrupt.js
baseball.js	divine.js
build.log	dull.js
football.js	grand.js
golf.js	lame.js
running.js	

For each pair of files, what percentage of lines are found in both?

	corrupt.js	divine.js	dull.js	grand.js	lame.js
archery.js					
baseball.js					
build.log					
football.js					
golf.js					
running.js					

Matrix of similarity percentages

How rename detection works

Crux of the problem

Rename detection is $O(M * N)$, where M and N are **huge**.

$\{M, N\} \sim O(\text{combined line count of potential rename } \{\text{sources}, \text{targets}\})$

Outline

- 1 "Merge machinery"
- 2 Merging and renames background
- 3 Strategies to improve rename performance
 - Exact renames
 - Partial capitulation
 - Dimensionality Reduction
 - Remembering previous work
 - But wait, there's more!
- 4 Results

Optimization 1: Don't redo work

Don't look for a better than perfect match.

Optimization 1: Don't redo work

Don't look for a better than perfect match.

Exact renames

Detecting renames

```
void detect_renames_and_copies(...)
{
    ...
    exact_count = find_different_name_same_hash();
    /* Keep all the source files as options for copies! */
    for (dest_path in potential_rename_targets) {
        if (already_paired(dest_path)) continue;
        for (source_path in potential_rename_sources) {
            if (!DETECT_COPIES ||
                already_paired(source_path))
                continue;
            compute_similarity();
        }
    }
    ...
}
```

Exact renames

Detecting renames

```
void detect_renames_and_copies(...)
{
    ...
    exact_count = find_different_name_same_hash();
    /* Keep all the source files as options for copies! */
    for (dest_path in potential_rename_targets) {
        if (already_paired(dest_path)) continue;
        for (source_path in potential_rename_sources) {
            if (!DETECT_COPIES ||
                already_paired(source_path))
                continue;
            compute_similarity();
        }
    }
    ...
}
```

Exact renames

Detecting renames

```
void detect_renames_and_copies(...)
{
    ...
    exact_count = find_different_name_same_hash();
    /* Keep all the source files as options for copies! */
    for (dest_path in potential_rename_targets) {
        if (already_paired(dest_path)) continue;
        for (source_path in potential_rename_sources) {
            if (!DETECT_COPIES &&
                already_paired(source_path))
                continue;
            compute_similarity();
        }
    }
    ...
}
```

Exact renames

Detecting renames

```
void detect_renames_and_copies(...)
{
    ...
    exact_count = find_different_name_same_hash();

    for (dest_path in potential_rename_targets) {
        if (already_paired(dest_path)) continue;
        for (source_path in potential_rename_sources) {
            if (!DETECT_COPIES &&
                already_paired(source_path))
                continue;
            compute_similarity();
        }
    }
    ...
}
```

Optimization 2: Don't do unnecessary work

If you can get the same answer without an expensive computation, skip the expensive computation.

Optimization 2: Don't do unnecessary work

If you can get the same answer without an expensive computation, skip the expensive computation.

Partial capitulation

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5caff01d	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5caff01d
Side2:	c01055a1
Merged:	0b57ac1e

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5caff01d	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5caff01d
Side2:	c01055a1
Merged:	0b57ac1e

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: modify/delete conflict
- viking.c: totally new file
- no textual merging

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	0b57ac1e

Then:

- buccaneer.c: removed
- viking.c: contains merged content

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **modify/delete conflict**
- viking.c: totally new file
- **no textual merging**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	???????

Then:

- buccaneer.c: removed
- viking.c: **contains merged content**

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **modify/delete conflict**
- viking.c: totally new file
- **no textual merging**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	???????

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **modify/delete conflict**
- viking.c: totally new file
- **no textual merging**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

EITHER

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

OR

```
Changes to be committed:
  deleted:   buccaneer.c
Unmerged paths:
  both modified:   viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **modify/delete conflict**
- viking.c: totally new file
- **no textual merging**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **deleted as expected**
- viking.c: totally new file
- **no textual merging**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```


Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **deleted as expected**
- viking.c: totally new file
- no textual merging **needed**

As reported by git status:

```
Changes to be committed:
  new file:   viking.c
Unmerged paths:
  deleted by them:   buccaneer.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **deleted as expected**
- viking.c: totally new file
- no textual merging **needed**

As reported by git status:

```
Changes to be committed:
  renamed:    buccaneer.c -> viking.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

```
Changes to be committed:
  renamed:    buccaneer.c -> viking.c
```

Partial capitulation

If a rename is not detected for the merge:

	buccaneer.c	viking.c
Base:	5eac0a57	00000000
Side1:	5eac0a57	00000000
Side2:	00000000	c01055a1

Then:

- buccaneer.c: **deleted as expected**
- viking.c: totally new file
- no textual merging **needed**

As reported by git status:

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

If we detect renames on each side of history:

	buccaneer.c ⇒ viking.c
Base:	5eac0a57
Side1:	5eac0a57
Side2:	c01055a1
Merged:	c01055a1

Then:

- buccaneer.c: removed
- viking.c: **no content merge was needed**

As reported by git status:

```
Changes to be committed:
  renamed:   buccaneer.c -> viking.c
```

Same results whether or not rename is detected by merge machinery.

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history.

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history.

Possible problems:

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history.

Possible problems:

- causes issues for directory rename detection

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history and parent directory of source file exists on *same* side of history.

Possible problems:

- causes issues for directory rename detection

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- Different conflict-related files in the working copy

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- Different conflict-related files in the working copy
- Different conflict-related entries in the index

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- Different conflict-related files in the working copy
- Different conflict-related entries in the index
- Different stdout; reports e.g. `CONFLICT (add/add)` instead of `CONFLICT (rename/add)`

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- Different conflict-related files in the working copy
- Different conflict-related entries in the index
- Different stdout; reports e.g. `CONFLICT (add/add)` instead of `CONFLICT (rename/add)`

After unifying file collision conflict handling...

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- ~~Different conflict related files in the working copy~~
- ~~Different conflict related entries in the index~~
- Different stdout; reports e.g. `CONFLICT (add/add)` instead of `CONFLICT (rename/add)`

After unifying file collision conflict handling...

Partial capitulation – Caveats?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history **and** parent directory of source file exists on *same* side of history.

Possible problems:

- ~~causes issues for directory rename detection~~
- rename/add conflict looks like add/add
- rename/rename(2to1) conflict looks like rename/add or add/add

“Mis-detected” conflict types:

- ~~Different conflict related files in the working copy~~
- ~~Different conflict related entries in the index~~
- Different stdout; reports e.g. `CONFLICT (add/add)` instead of `CONFLICT (rename/add)`

After unifying file collision conflict handling...**stdout is only difference.**

Partial capitulation – micro or mega optimization?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history and parent directory of source file exists on *same* side of history.

Partial capitulation – micro or mega optimization?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history and parent directory of source file exists on *same* side of history.

How much does this new strategy help?

Partial capitulation – micro or mega optimization?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history and parent directory of source file exists on *same* side of history.

How much does this new strategy help?

A Common Case

$$O(M * N) \rightarrow$$

Partial capitulation – micro or mega optimization?

New Strategy

Exclude potential source from rename detection **if** it is unmodified by *other* side of history and parent directory of source file exists on *same* side of history.

How much does this new strategy help?

A Common Case

$$O(M * N) \rightarrow O(\emptyset * N)$$

Optimization 3: Fudge “unnecessary”

Only do part of the work and accept slightly different results if there are huge cost savings.

Optimization 3: Fudge “unnecessary”

Only do part of the work and accept slightly different results if there are huge cost savings.

Dimensionality Reduction

Fun fact

Over 75% of renames in the linux kernel repository do not change the basename of the file, just the directory in which it is found.

Dimensionality Reduction

Detecting renames...

Files in Base	Files in given side

For each pair of files, what percentage of lines are found in both?

--	--	--	--

Dimensionality Reduction

Detecting renames...

Files in Base	Files in given side
document.html	build.log
src/ blue.css	document.html
src/ brown.css	source/ blue.css
src/ green.css	source/ brown.css
src/ red.css	source/ green.css
	source/ orange.css
	source/ purple.css
	source/ red.css

For each pair of files, what percentage of lines are found in both?

Dimensionality Reduction

Detecting renames...

Files in Base	Files in given side
document.html	build.log
src/blue.css	document.html
src/brown.css	source/blue.css
src/green.css	source/brown.css
src/red.css	source/green.css
	source/orange.css
	source/purple.css
	source/red.css

For each pair of files, what percentage of lines are found in both?

	src/blue.css	src/brown.css	src/green.css	src/red.css
build.log				
source/blue.css				
source/brown.css				
source/green.css				
source/orange.css				
source/purple.css				
source/red.css				

Dimensionality Reduction

Detecting renames...

Files in Base	Files in given side
document.html	build.log
src/blue.css	document.html
src/brown.css	source/blue.css
src/green.css	source/brown.css
src/red.css	source/green.css
	source/orange.css
	source/purple.css
	source/red.css

For each pair of files, what percentage of lines are found in both?

	src/blue.css	src/brown.css	src/green.css	src/red.css
build.log				
source/blue.css				
source/brown.css				
source/green.css				
source/orange.css				
source/purple.css				
source/red.css				

Matrix of similarity percentages

Dimensionality Reduction

Detecting renames...

Files in Base	Files in given side
document.html	build.log
src/blue.css	document.html
src/brown.css	source/blue.css
src/green.css	source/brown.css
src/red.css	source/green.css
	source/orange.css
	source/purple.css
	source/red.css

For each pair of files, what percentage of lines are found in both?

	src/blue.css	src/brown.css	src/green.css	src/red.css
build.log				
source/blue.css				
source/brown.css				
source/green.css				
source/orange.css				
source/purple.css				
source/red.css				

Sparse similarity percentages

Dimensionality Reduction

Improvement

$$O(M * N) \rightarrow O((M-B)*(N-B) \quad)$$

Dimensionality Reduction

Improvement

$$O(M * N) \rightarrow O((M-B)*(N-B) + B)$$

Dimensionality Reduction

Improvement

$$O(M * N) \rightarrow O((M-B)*(N-B) + B)$$

If enough matching basenames...

$$O(M * N) \rightarrow O(\text{minimum}(M, N))$$

Optimization 4: Don't redo unnecessary work


When repeatedly merging, re-use previous rename detection results.

Optimization 4: Don't redo unnecessary work

When repeatedly merging, re-use previous rename detection results.


Remembering previous work

Cherry-picking C..feature




Remembering previous work

Picking F




Remembering previous work

Picking F



Remembering previous work


Picking F



	buccaneer.c	viking.c
C :	c0a575	000000
E :	000000	b1ade5
F :	befall	000000

Remembering previous work


Picking F



$$\left\{ \begin{array}{ll} \text{buccaneer.c} & \text{viking.c} \\ C: & \text{c0a575} \quad 000000 \\ E: & 000000 \quad \text{blade5} \\ F: & \text{befa11} \quad 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ C: \text{c0a575} \\ E: \text{blade5} \\ F: \text{befa11} \end{array} \right\}$$

Remembering previous work


Picking F



$$\left\{ \begin{array}{cc} \text{buccaneer.c} & \text{viking.c} \\ C: & \text{c0a575} \quad 000000 \\ E: & 000000 \quad \text{blade5} \\ F: & \text{befa11} \quad 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ C: \text{c0a575} \\ E: \text{blade5} \\ F: \text{befa11} \\ F^*: 1007\text{ed} \end{array} \right\}$$

Remembering previous work

Picked F



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

$$\left\{ \begin{array}{ll} \text{buccaneer.c} & \text{viking.c} \\ C: & \text{c0a575} \quad 000000 \\ E: & 000000 \quad \text{b1ade5} \\ F: & \text{befa11} \quad 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ C: \text{c0a575} \\ E: \text{b1ade5} \\ F: \text{befa11} \\ F*: 1007ed \end{array} \right\}$$

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575


E: b1ade5

F: befa11

*F**: 1007ed

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575


E: b1ade5

F: befa11

*F**: 1007ed

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

$$\left\{ \begin{array}{ll} \text{buccaneer.c} & \text{viking.c} \\ F: & \text{befa11} \quad 000000 \\ F*: & 000000 \quad 1007ed \\ G: & \text{a70115} \quad 000000 \end{array} \right\}$$

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

	buccaneer.c	viking.c
<i>F</i> :	befa11	000000
<i>F*</i> :	000000	1007ed
<i>G</i> :	a70115	000000

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

	buccaneer.c	viking.c
<i>F</i> :	befa11	000000
<i>F*</i> :	000000	1007ed
<i>G</i> :	a70115	000000

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

$$\left\{ \begin{array}{cc} & \text{buccaneer.c} & \text{viking.c} \\ F: & \text{befa11} & 000000 \\ F*: & 000000 & 1007ed \\ G: & a70115 & 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ F: \text{befa11} \\ F*: 1007ed \\ G: a70115 \end{array} \right\}$$

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C: c0a575

E: b1ade5


F: befa11

*F**: 1007ed

$$\left\{ \begin{array}{ll} \text{buccaneer.c} & \text{viking.c} \\ F: & \text{befa11} \quad 000000 \\ F^*: & 000000 \quad 1007ed \\ G: & a70115 \quad 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ F: \text{ befa11} \\ F^*: 1007ed \\ G: a70115 \\ G^*: \text{ fabled} \end{array} \right\}$$

Remembering previous work

Picking G



buccaneer.c \Rightarrow viking.c

C : c0a575

E : b1ade5

F : befa11

F^* : 1007ed

$$\left\{ \begin{array}{ll} \text{buccaneer.c} & \text{viking.c} \\ F: & \text{befa11} \quad 000000 \\ F^*: & 000000 \quad 1007ed \\ G: & a70115 \quad 000000 \end{array} \right\} \rightarrow \left\{ \begin{array}{l} \text{buccaneer.c} \Rightarrow \text{viking.c} \\ F: \text{befa11} \\ F^*: 1007ed \\ G: a70115 \\ G^*: \text{fab1ed} \end{array} \right\}$$

But wait, there's more!

- Avoid accidentally quadratic behavior
- Restructure to eliminate quasi-quadratic index insertion and removal
- Fewer tree traversals
- Extend “partial capitulation” ideas from *file* renames to *directory* renames
- Avoid updating the index or working tree if not needed
 - Helps with new sparse-checkout command
 - Accelerates rebases and cherry-picks
 - Avoids unnecessary recompilation after a rebase
- ...and a few other minor improvements

But wait, there's more!

- Avoid accidentally quadratic behavior
- Restructure to eliminate quasi-quadratic index insertion and removal
- Fewer tree traversals
- Extend “partial capitulation” ideas from *file* renames to *directory* renames
- Avoid updating the index or working tree if not needed
 - Helps with new **sparse-checkout** command
 - Accelerates rebases and cherry-picks
 - Avoids unnecessary recompilation after a rebase
- ...and a few other minor improvements

But wait, there's more!

- Avoid accidentally quadratic behavior
- Restructure to eliminate quasi-quadratic index insertion and removal
- Fewer tree traversals
- Extend “partial capitulation” ideas from *file* renames to *directory* renames
- Avoid updating the index or working tree if not needed
 - Helps with new **sparse-checkout** command
 - Accelerates rebases and cherry-picks
 - **Avoids unnecessary recompilation** after a rebase
- ...and a few other minor improvements

Outline

- 1 “Merge machinery”
- 2 Merging and renames background
- 3 Strategies to improve rename performance
- 4 Results

Testcase

- Linux kernel
 - Rebase or cherry-pick hwmon-updates (35 patches) from 5.5 \rightarrow 5.4
 - Very few renames involved; only takes 50% of execution time
 - Speedup factor of 3 (optimized more things than renames)
 - What if we checkout 5.4, and rename `drivers/` \Rightarrow `pilots/`, and then rebase or cherry-pick those 35 patches on top ?

Testcase

- Linux kernel
- Rebase or cherry-pick hwmon-updates (35 patches) from 5.5 \rightarrow 5.4
- Very few renames involved; only takes 50% of execution time
- Speedup factor of 3 (optimized more things than renames)
- What if we checkout 5.4, and rename `drivers/` \Rightarrow `pilots/`, and then rebase or cherry-pick those 35 patches on top ?

Testcase

- Linux kernel
- Rebase or cherry-pick hwmon-updates (35 patches) from 5.5 \rightarrow 5.4
- Very few renames involved; only takes 50% of execution time
- Speedup factor of 3 (optimized more things than renames)
- What if we checkout 5.4, and rename `drivers/` \Rightarrow `pilots/`, and then rebase or cherry-pick those 35 patches on top ?

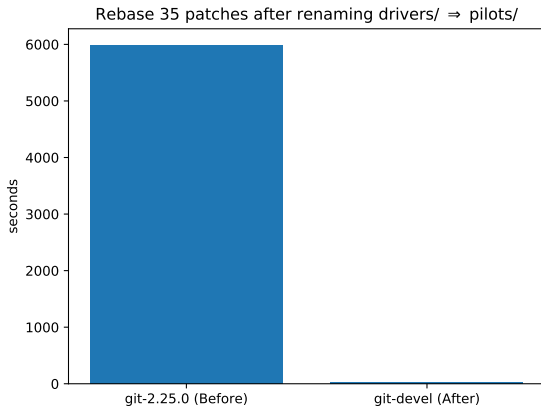
Testcase

- Linux kernel
- Rebase or cherry-pick hwmon-updates (35 patches) from 5.5 \rightarrow 5.4
- Very few renames involved; only takes 50% of execution time
- Speedup factor of 3 (optimized more things than renames)
- What if we checkout 5.4, and rename `drivers/` \Rightarrow `pilots/`, and then rebase or cherry-pick those 35 patches on top ?

Testcase

- Linux kernel
- Rebase or cherry-pick hwmon-updates (35 patches) from 5.5 \rightarrow 5.4
- Very few renames involved; only takes 50% of execution time
- Speedup factor of 3 (optimized more things than renames)
- What if we checkout 5.4, and rename `drivers/` \Rightarrow `pilots/`, and then rebase or cherry-pick those 35 patches on top ?

Results



Results

Reproduce these numbers:

<https://github.com/newren/git/blob/git-merge-2020-demo/README.md>

The journey, redux

Issues starting the journey

- cherry-pick would fail to detect renames and fail to notify about needed `merge.renameLimit`
- cherry-pick would ignore `merge.renameLimit > 32767`
- if directory renames involved, files would be left in wrong directory
- people wrote custom purpose scripts to cherry-pick things
- after fixing `merge.renameLimit`, cherry-picking small patches would take more than 9 minutes.

The journey, redux

Issues starting the journey

- cherry-pick would fail to detect renames and fail to notify about needed `merge.renameLimit`
- cherry-pick would ignore `merge.renameLimit > 32767`
- if directory renames involved, files would be left in wrong directory
- people wrote custom purpose scripts to cherry-pick things
- after fixing `merge.renameLimit`, cherry-picking small patches would take more than 9 minutes.

When I told folks a few years ago that "You don't need these special scripts to cherry pick things; just set `merge.renameLimit` to something higher," they responded that `merge.renameLimit` didn't work.

The journey, redux

Issues starting the journey

- cherry-pick would fail to detect renames and fail to notify about needed merge.renameLimit
- cherry-pick would ignore merge.renameLimit > 32767
- if directory renames involved, files would be left in wrong directory
- people wrote custom purpose scripts to cherry-pick things
- after fixing merge.renameLimit, cherry-picking small patches would take more than 9 minutes.

When I told folks a few years ago that "You don't need these special scripts to cherry pick things; just set merge.renameLimit to something higher," they responded that merge.renameLimit didn't work.

I didn't believe them.

The journey, redux

Issues starting the journey

- cherry-pick would fail to detect renames and fail to notify about needed merge.renameLimit
- cherry-pick would ignore merge.renameLimit > 32767
- if directory renames involved, files would be left in wrong directory
- people wrote custom purpose scripts to cherry-pick things
- after fixing merge.renameLimit, cherry-picking small patches would take more than 9 minutes.

When I told folks a few years ago that "You don't need these special scripts to cherry pick things; just set merge.renameLimit to something higher," they responded that merge.renameLimit didn't work.

I didn't believe them.

My efforts in this area, including this performance work, represent my attempt to continue to not believe them. :-)